

Revision of the paper

“Multiscale inference and long-run variance estimation in nonparametric regression with time series errors”

First of all, we would like to thank the editor, the associate editor and the reviewers for their many comments and suggestions which were very helpful in improving the paper. In the revision, we have addressed all comments and have rewritten the paper accordingly. Please find our point-by-point responses below. Since the revised paper includes additional material as requested by the referees (additional simulations, a second application example, ...), it is a bit longer than the original submission. In particular, it has grown from 34 to 37 pages in our layout. However, we are of course happy and willing to reduce the length of the paper if this is needed. Before we reply to the specific comments of the referees, we summarize the major changes in the revision.

Generalization of the theoretical results. We have extended the theoretical results as suggested by Referee 1:

- (i) We have derived the following consistency result in addition to Proposition 3.3:
Let the significance level $\alpha = \alpha_T \in (0, 1)$ depend on the sample size T . If $\alpha_T \rightarrow 0$, then $\mathbb{P}(E_T^\ell) \rightarrow 1$ for $\ell \in \{\pm, +, -\}$. This result is stated as Corollary 3.1 on p.13 of the revised paper.
- (ii) We have generalized our estimator of the long-run error variance. The estimation procedure is shown to be valid not only for $\text{AR}(p)$ processes of known finite order p but for any stationary error process $\{\varepsilon_t\}$ with an $\text{AR}(\infty)$ representation.

Comparison to SiZer. In the new Section 3.4, we give a clear account of the main contributions and innovations of our paper relative to the SiZer approach. Moreover, we have thoroughly revised the simulation study in Section 5 which compares our multiscale test with SiZer, taking into account the suggestions of Referee 2. In what follows, we summarize the main points of the new Section 3.4.

Informally speaking, both our approach and SiZer for dependent data (dependent SiZer for short) are methods to test for local increases/decreases of a nonparametric trend function m . The formal problem is to test the hypothesis

$$H_0(u, h) : \text{The trend } m \text{ is constant on the time interval } [u - h, u + h]$$

simultaneously for a large number of different time intervals $[u - h, u + h]$, in particular, for all intervals with $u \in I_T$ and $h \in H_T$, where I_T is the set of locations and H_T the set of bandwidths/scales h under consideration.

Let $\widehat{s}_T(u, h)$ be the SiZer statistic to test $H_0(u, h)$, which corresponds to the statistic $\widehat{\psi}_T(u, h)/\widehat{\sigma}$ in our approach and which is properly defined in Section 3.4. There are two versions of dependent SiZer:

- (a) The global version aggregates the individual statistics $\widehat{s}_T(u, h)$ into the overall statistic $\widehat{S}_T = \max_{h \in H_T} \widehat{S}_T(h)$, where $\widehat{S}_T(h) = \max_{u \in I_T} |\widehat{s}_T(u, h)|$. The statistic \widehat{S}_T is the counterpart to the multiscale statistic $\widehat{\Psi}_T$ of our approach.
- (b) The row-wise SiZer version considers each scale $h \in H_T$ separately. In particular, for each bandwidth $h \in H_T$, a test is carried out based on the statistic $\widehat{S}_T(h)$.

In practice, SiZer is commonly implemented in its row-wise form. The main reason is that it has more power than the global version by construction. However, this gain of power comes at a cost: Row-wise SiZer carries out a test *separately* for each scale $h \in H_T$, thus ignoring the simultaneous test problem across scales h . Hence, it is not a rigorous level- α -test of the overall null hypothesis H_0 . For this reason, we focus on global SiZer in what follows.

Even though related, our methods and theory are markedly different from those of the SiZer approach:

- (i) Theory for SiZer is derived under the following assumption on the set of bandwidths H_T : $H_T \subseteq H$ for all T , where H is a compact subset of $(0, \infty)$. As already pointed out in Chaudhuri and Marron (2000), this is a quite severe restriction: Only bandwidths h are taken into account that remain bounded away from zero as the sample size T grows. Bandwidths h that converge to zero as T increases are excluded. As Chaudhuri and Marron (2000) put it (on p.420):

Note that all the weak convergence results in this section have been established under the assumption that both of H and I are fixed compact subintervals of $(0, \infty)$ and $(-\infty, \infty)$ respectively. Compactness of the set $H \times I$ enables us to exploit standard results on weak convergence of a sequence of probability measures on a space of continuous functions defined on a common compact metric space. However, conventional asymptotics for nonparametric curve estimates allows the smoothing parameter h to shrink with growing sample size. There frequently one assumes that h_n is of the order $n^{-\gamma}$ for some appropriate choice of $0 < \gamma < 1$ so that the estimate $\widehat{f}_{h_n}(x)$ converges to the “true function” $f(x)$ at an “optimal rate”. This makes one wonder about the asymptotic behavior of the empirical scale space surface when h varies in $H_n = [an^{-\gamma}, b]$, where $a, b > 0$ are fixed constants. Extension of our weak convergence results along that direction will be quite interesting, and we leave it as a challenging open problem here.

The theory of our paper allows to deal with this problem. In particular, it allows to simultaneously consider scales h that remain bounded away from zero and scales $h = h_T$ that converge to zero at various different rates $T^{-\gamma}$. To achieve

this, we come up with a proof strategy which is very different from that in the SiZer literature: As proven in Chaudhuri and Marron (2000) for the i.i.d. case and in Park et al. (2009) for the dependent data case, \widehat{S}_T weakly converges to some limit S under the overall null hypothesis H_0 . This is the central technical result on which the theoretical properties of SiZer are based. In contrast to this, our proof strategy (which combines strong approximation theory with anti-concentration bounds as outlined in Section 3.3 of the paper) does not even require the statistic $\widehat{\Psi}_T$ to have a weak limit and is thus not restricted by the limitations of classic weak convergence theory.

- (ii) There are different ways to combine the test statistics $\widehat{S}_T(h) = \max_{u \in I_T} |\widehat{s}_T(u, h)|$ for different scales $h \in H_T$. One way is to take their maximum, which leads to the SiZer statistic $\widehat{S}_T = \max_{h \in H_T} \widehat{S}_T(h)$. We could proceed analogously and consider the statistic $\widehat{\Psi}_{T, \text{uncorrected}} = \max_{h \in H_T} \widehat{\Psi}_T(h) = \max_{(u, h) \in I_T \times H_T} |\widehat{\psi}_T(u, h)/\widehat{\sigma}|$. However, as argued in Dümbgen and Spokoiny (2001) and as discussed in Section 3.1 of our paper, this aggregation scheme is not optimal when the set H_T contains scales h of many different rates. Following the lead of Dümbgen and Spokoiny (2001), we consider the test statistic $\widehat{\Psi}_T = \max_{(u, h) \in I_T \times H_T} \{|\widehat{\psi}_T(u, h)/\widehat{\sigma}| - \lambda(h)\}$ with the additive correction terms $\lambda(h)$. Hence, even though related, our multiscale test statistic $\widehat{\Psi}_T$ differs from the SiZer statistic \widehat{S}_T in important ways.
- (iii) The main complication in carrying out both our multiscale test and SiZer is to determine the critical values, that is, the quantiles of the test statistics $\widehat{\Psi}_T$ and \widehat{S}_T under H_0 . In order to approximate the quantiles, we proceed quite differently than in the SiZer literature. The quantiles of the SiZer statistic \widehat{S}_T can be approximated by those of the weak limit S . Usually, however, the quantiles of S cannot be determined analytically but have to be approximated themselves (e.g. by the bootstrap procedures of Chaudhuri and Marron (1999, 2000)). Alternatively, the quantiles of \widehat{S}_T can be approximated by procedures based on extreme value theory (as proposed in Hannig and Marron (2006) and Park et al. (2009)). In our approach, the quantiles of $\widehat{\Psi}_T$ under H_0 are approximated by those of a suitably constructed Gaussian analogue of $\widehat{\Psi}_T$. It is far from obvious that this Gaussian approximation is valid when the data are dependent. To see this, deep strong approximation theory for dependent data (as derived in Berkes et al. (2014)) is needed. It is important to note that our Gaussian approximation procedure is not the same as the bootstrap procedures proposed in Chaudhuri and Marron (1999, 2000). Both procedures can of course be regarded as resampling methods. However, the resampling is done in a quite different way in our case.

We hope the above points make clear that the methodological and theoretical contributions of our paper are quite substantial relative to the SiZer methodology.

Simulations and application examples. We have thoroughly revised the simulation study of Section 5. In order to take into account the many suggestions of Referees 1 and 2, we have completely re-designed the simulation exercises which analyse the finite sample properties of our multiscale test and dependent SiZer. These were formerly presented in Sections 5.1 and 5.2 and are combined in the new Section 5.1 of the revised paper. Among other things, we consider additional simulation designs (with different trend functions and AR error terms) which are chosen according to the suggestions of the referees. To keep the simulation study to a reasonable length, we have removed some of the simulation setups from the previous version of the paper, in particular, the setup with AR(2) errors that mimics the situation in the application example. Finally, we have added a second application example to global temperature data as requested by Referee 1.

The revised R-code to run the simulation exercises and the application examples has been submitted as part of the revision. When revising the code, we found a minor mistake in the computation of the long-run error variance estimator $\hat{\sigma}^2$. This mistake, however, only concerned the case of AR(p) errors with $p > 1$. Hence, it did not affect the simulations on the long-run error variance estimator in Section 5.2, where $p = 1$ throughout. Only the long-run variance estimate $\hat{\sigma}^2$ in the application to the Central England temperature record (where $p = 2$) changed a bit. The revised application example can be found in the new Section 6.1. As can be seen, the corrected estimate $\hat{\sigma}^2 = 0.737$ is quite close to the old one $\hat{\sigma}^2 = 0.749$ and the test results are qualitatively the same as before. The R-code to produce our simulation and application results is accompanied by a README file which explains the main structure of the code and how to run it.

Reply to Referee 1

Thank you very much for the constructive and helpful comments. In our revision, we have addressed all of them. Please see our replies to your comments below.

(1) *A consistency result of Proposition 3.3.*

I believe that the following type of result can be obtained: $\mathbb{P}(E_T^\ell) \rightarrow 1$. Theorem 3.1 is for testing purpose. In certain application one might be interested in such consistency result. Basically one needs to study the behavior of $q_T(\alpha)$ when $\alpha \rightarrow 0$.

Under our regularity conditions, it can indeed be proven that $\mathbb{P}(E_T^\ell) \rightarrow 1$ as $\alpha = \alpha_T \rightarrow 0$. We have added this consistency result as Corollary 3.1 to the paper. The proof is provided in the Supplementary Material.

(2) *Estimation of long run variance using autoregressive processes.*

The authors considered estimating σ^2 using AR processes. A limitation is that the order p is fixed and finite. It appears that the latter limitation can be relaxed. For a stationary process ε_t (not necessarily linear), one can fit an AR process with large p ,

$$\varepsilon_t = \sum_{j=1}^p a_j \varepsilon_{t-j} + \eta_t,$$

properties of fitted $\hat{a}_1, \dots, \hat{a}_p$ can be obtained from the results in the following papers: Wu and Pourahmadi (2009) and Xiao and Wu (2012). A similar version of the authors' estimate (4.14) can be used. Rate of convergence (cf. Proposition 4.1) can be derived with rate $T^{-1/2}$ therein possibly replaced by a larger term of the form T^{-c} with $c < 1/2$.

Many thanks for this interesting suggestion. We have generalized our procedure to estimate the long-run variance σ^2 along the lines suggested by you: Rather than considering $\text{AR}(p)$ processes of known finite order p , we consider the much more general class of $\text{AR}(\infty)$ processes, which nests AR processes of any finite order as a special case. More specifically, we assume the error process $\{\varepsilon_t\}$ to have the form

$$\varepsilon_t = \sum_{j=1}^{\infty} a_j \varepsilon_{t-j} + \eta_t, \tag{*}$$

where a_1, a_2, a_3, \dots are unknown parameters and η_t are i.i.d. innovations (which fulfill certain regularity conditions as detailed in the revised Section 4). As far as we can see, this is the most general class of error processes to which we can extend our methods: For our theory to work, we require the process $\{\varepsilon_t\}$ to have at least an $\text{AR}(\infty)$ representation because otherwise we do not get suitable Yule-Walker equations that we can exploit.

In the paper, we reformulate (*) as follows: We assume that $\{\varepsilon_t\}$ has the form

$$\varepsilon_t = \sum_{j=1}^{p^*} a_j \varepsilon_{t-j} + \eta_t, \quad (**)$$

where $p^* \in \mathbb{N} \cup \{\infty\}$ is the true (unknown) AR order of $\{\varepsilon_t\}$ which may be finite or infinite. In order to generalize our theory to the case that $\{\varepsilon_t\}$ has an $\text{AR}(p^*)$ representation of the form (**), we fit $\text{AR}(p)$ type models to the data, where we distinguish between the following two cases:

- (A) We do not know the precise AR order p^* but we know an upper bound $p \in \mathbb{N}$ with $p \geq p^*$.
- (B) We neither know p^* nor an upper bound on it. In this case, we let $p = p_T \rightarrow \infty$ as $T \rightarrow \infty$.

Whereas p is fixed in case (A), we fit $\text{AR}(p)$ type processes of growing order $p = p_T$ to the data in case (B). We thus approximate the $\text{AR}(p^*)$ process $\{\varepsilon_t\}$ by a sequence of $\text{AR}(p)$ processes whose order $p = p_T$ goes to infinity. This approach is somewhat simpler but related to the banding techniques developed in Wu and Pourahmadi (2009) and Xiao and Wu (2012).

The convergence rates of our estimators are derived in Proposition 4.1 for both cases (A) and (B). Whereas the estimators are \sqrt{T} -consistent in case (A), the convergence rates are a bit slower in case (B) as you have already conjectured in your comment. Another technical detail which is different in cases (A) and (B) is the following: The second-step estimator $\hat{\mathbf{a}}_r$ of the AR parameters $\mathbf{a} = (a_1, \dots, a_p)^\top$ depends on the tuning parameter r , which can be chosen as any fixed natural number in case (A). In case (B), in contrast, r is required to be slightly larger than p , in particular, $r \geq (1 + \delta)p$ for some arbitrarily small but fixed $\delta > 0$. The technical reason for this is as follows: Let $\mathbf{\Gamma}_r = (\gamma_r(i - j) : 1 \leq i, j \leq p)$ with $\gamma_r(\ell) = \text{Cov}(\Delta_r \varepsilon_t, \Delta_r \varepsilon_{t-\ell})$ be the autocovariance matrix of the r -th differences $\Delta_r \varepsilon_t = \varepsilon_t - \varepsilon_{t-r}$. Notably, $\mathbf{\Gamma}_r$ is of growing dimension $p = p_T$ in case (B). To derive our theoretical results, we need to show the following:

- (+) $\|\mathbf{\Gamma}_r^{-1}\|_2 \leq C$ for sufficiently large T with C independent of p , r and T .

A standard strategy to prove statements such as (+) is to bound the eigenvalues of $\mathbf{\Gamma}_r$ by invoking results on Toeplitz matrices (see e.g. Section 5.2 in Grenander and Szegö (1958) or Proposition 4.5.3 in Brockwell and Davis (1991)). However, these results yield (+) only if the spectral density of $\{\Delta_r \varepsilon_t\}$ is bounded away from zero. Unfortunately, this is *not* the case: $\{\Delta_r \varepsilon_t\}$ is an $\text{ARMA}(p^*, r)$ process with a unit root in the MA polynomial, implying that its spectral density takes the value 0.

We thus had to prove (+) by a different strategy, which is given in Lemma S.7 of the Supplement. For this strategy to work, we require that $r \geq (1 + \delta)p$.

Our revised and extended methods to estimate the AR parameters and the long-run variance of the error process $\{\varepsilon_t\}$ can be found in Section 4 of the paper. The proof of Proposition 4.1 is provided in Section S.2 of the Supplement. Please also note that we have removed Section 4.1 (which discusses long-run variance estimation for general weakly dependent processes) from the paper as requested by Referee 2.

(3) *Real data application.*

The authors analyzed the yearly mean Central England temperature data. It will be interesting to apply their approach to the global temperature data. In the paper Wu et al. (2001), an increasing trend function is fitted. It will be important to know in which period the sequence is increasing/decreasing.

We have added the application to global temperature data as a second empirical example. In particular, we have used exactly the same data as in Wu et al. (2001) in order to be able to compare our test results with theirs. Please see the new Section 6.2 for the details.

(4) *Simulation study.*

In the simulation study, the authors considered AR(1) processes with relatively weak dependence: $a_1 \in \{-0.5, -0.25, 0.25, 0.5\}$. One should consider the stronger positive/negative dependence case with $a = \pm 0.9$ (say). How does the strength of dependence affect the performance of the procedure?

We have added the AR(1) case with $a = \pm 0.9$ to our revised simulation study in Section 5.1. In particular, we have carried out additional size simulations for the case that $a = \pm 0.9$. The results are reported in Table 2 on p.22 and can be summarized as follows: The size of our multiscale test gives a decent approximation to the nominal target α for sample sizes $T \geq 1000$. However, for the smaller sample sizes $T = 250$ and $T = 500$, there are substantial size distortions. Hence, when the error terms are strongly autocorrelated, our test has good size properties only for sufficiently large sample sizes. This is indeed what we have expected: Statistical inference in the presence of strongly autocorrelated data is a hard problem in general and satisfactory results can only be expected for reasonably large sample sizes.

Reply to Referee 2

Thank you very much for the constructive and useful suggestions. In our revision, we have addressed all of them. In particular, we have thoroughly revised the simulation study which compares our multiscale test with SiZer according to your suggestions. Here are our point-by-point responses to your comments.

- (1) *Section 3.2: The authors recommend computing the quantiles for the independent Gaussian case by simulation. This suggestion is already in the original SiZer paper (Chaudhuri and Marron, 1999). However in the late 1990s computing power was not sufficient to make this suggestion feasible. This led to the use of approximation such as in Hannig and Marron (2006). I would like to ask how does the simulation based quantile compare to the approximation in Hannig and Marron (2006).*

If we understand you correctly, your question is the following: How does our simulation-based test procedure compare (empirically) to the approximation-based SiZer version of Hannig and Marron (2006) and Park et al. (2009)? We hope this is indeed what you mean by the question “how does the simulation based quantile compare to the approximation in Hannig and Marron (2006)”.

We have done the following to answer this question: In the revised simulation study of Section 5.1, we consider a row-wise version \mathcal{T}_{RW} of our multiscale test, which is the counterpart to row-wise dependent SiZer developed in Park et al. (2004), Rondonotti et al. (2007) and Park et al. (2009). A definition of the row-wise test \mathcal{T}_{RW} is given on the top of p.21. The quantiles of \mathcal{T}_{RW} under the null are approximated by our simulation-based procedure, whereas the quantiles of row-wise dependent SiZer are approximated by the method of Hannig and Marron (2006) and Park et al. (2009) which is based on extreme value theory. In our simulation study, we compare the size and power properties of the row-wise version \mathcal{T}_{RW} of our test and row-wise dependent SiZer. This shows how our simulation-based procedure compares to the approximation-based SiZer procedure of Hannig and Marron (2006) and Park et al. (2009).

Before moving on to your next comment, we would like to emphasize the following two points on our simulation-based test procedure and its relation to the simulation/bootstrap procedures proposed in the original SiZer paper (Chaudhuri and Marron, 1999):

- As already mentioned in the summary of the main differences between our multiscale approach and SiZer at the beginning of this letter, the Gaussian approximation procedure that we use for simulating the quantiles of the multiscale statistic under the null is not the same as the (empirical) bootstrap procedures

proposed in Chaudhuri and Marron (1999, 2000). Both procedures can be regarded as resampling methods. However, the resampling is done in a quite different way.

- It is far from obvious that our simulation-based procedure is theoretically valid and provides an adequate critical value such that our test has asymptotically the correct size under the null. One of the main theoretical contributions of our paper is to formally show that this is indeed the case.

- (2) *Page 12, line 15: What is random here? After a spending some time I believe that it is the Π_T but on first reading I thought E_T s are non-random. Please explain these various objects better.*

We have rewritten the text concerning the objects Π_T^ℓ and E_T^ℓ for $\ell \in \{\pm, +, -\}$. We have attempted to explain these objects in a clearer way and to clarify the question what is random here. (The collection of intervals Π_T^ℓ is indeed random.) Please see p.11/12 for the details. We hope the new exposition of the material makes things clearer.

- (3) *Page 18, line 52: Please remove the speculative statements about what can be shown unless you actually show it in this paper.*

We have removed the speculative statements following Proposition 4.1.

- (4) *Section 4.1: This section does not contain any truly new material and should be removed.*

We have removed Section 4.1 from the paper. Moreover, we have revised Section 4 on the estimation of the long-run error variance according to the suggestions of Referee 1. Please see our reply to comment (2) of Referee 1 for the details.

- (5) *Section 5.2: I understand that you are doing comparisons to SiZer out of the box. However, some of the comparison might not be quite fair. SiZer is adjusting multiplicity row-wise while the proposed method is attempting a global multiple control. What would happen if your \mathcal{G}_T only focused on one scale?*

We have thoroughly revised the comparison study with SiZer, attempting to give a fairer comparison of the methods. The revised study compares the following versions of our multiscale test and SiZer:

- our multiscale test \mathcal{T}_{MS} as defined in Section 3 of the paper
- an uncorrected version \mathcal{T}_{UC} of our multiscale test without the additive correction terms $\lambda(h)$
- a row-wise version \mathcal{T}_{RW} of our multiscale test

- row-wise dependent SiZer $\mathcal{T}_{\text{SiZer}}$ of Park et al. (2004), Rondonotti et al. (2007) and Park et al. (2009).

A brief description of the four test methods \mathcal{T}_{MS} , \mathcal{T}_{UC} , \mathcal{T}_{RW} and $\mathcal{T}_{\text{SiZer}}$ is given at the beginning of the simulation section on p.20/21. We have analyzed the size and power properties of the four methods, in particular, not only the global but also the row-wise size and power properties to be fairer with regard to SiZer. As already mentioned in our answer to your comment (1), the row-wise version \mathcal{T}_{RW} of our multiscale test focuses on one scale at a time and is thus the direct counterpart to $\mathcal{T}_{\text{SiZer}}$. Please see Section 5.1 for the full simulation study, which in particular shows how \mathcal{T}_{RW} and $\mathcal{T}_{\text{SiZer}}$ compare to each other. We hope you find the revised study more accurate.

- (6) *Page 25, line 1-26: I do not quite understand this figure. Would it be possible to rather reproduce the colorful SiZer figures that show the results of the test at various scales and locations? Also you should use several different signals. I believe that a single relatively large bump is not sufficient test bed. A good collection of signals can be found in Donoho and Johnstone (1995). Also, would Hannig et al. (2013) be helpful in comparing the results?*

We have removed the figure from the paper. In order to compare the power properties of our multiscale approach and SiZer, we perform the following simulation exercises in the revision:

- (a) As suggested by you, we consider different trend signals and produce the SiZer plots for both our approach and SiZer. We use the following signals: (i) the sine signal $m(u) = \sin(6\pi u)$ that was considered in Park et al. (2009) and (ii) the blocks signal from Donoho and Johnstone (1995) which was investigated in detail in Hannig and Marron (2006). In both cases, the error terms are modelled as an AR(1) process with the parameter $a_1 \in \{-0.5, 0.5\}$.
- (b) In addition to (a), we perform more or less conventional power comparisons for our approach and SiZer. To do so, we consider a simple trend signal, in particular, a bump signal as in the previous version of the paper. The considered bump signal is increasing (i.e. has a positive derivative) on the interval $I^+ = (0.45, 0.5)$, is decreasing on $I^- = (0.5, 0.55)$ and is constant elsewhere. Suppose we are interested in finding local increases in m . When performing our multiscale test and SiZer, we can distinguish between
 - a correct finding: the test finds an increase on some interval $I_{u,h} = [u - h, u + h]$ which intersects with I^+ (that is, the test finds an increase on some interval $I_{u,h}$ where m is indeed increasing).

- a spurious/false finding: the test finds an increase on some interval $I_{u,h} = [u - h, u + h]$ which does not intersect with I^+ (that is, the test finds an increase on some interval $I_{u,h}$ where m is not increasing).

For both our approach and SiZer, we define the following concepts:

- global power: the number of simulations where there is at least one correct finding divided by the total number of simulations
- spurious global power: the number of simulations where there is at least one spurious/false finding divided by the total number of simulations.

In addition, we define row-wise power and spurious row-wise power in an analogous way. We now proceed as follows: We simulate $S = 1000$ data samples and compute (global/row-wise) power as well as spurious (global/row-wise) power for the tests under consideration.

The simulation exercises from (b) can be found in Section 5.1.2. Those from (a) are in Section S.3 of the Supplementary Material. We have decided not to include them in the paper because of space considerations. (The figures with the SiZer plots are already two full pages by themselves.) However, if you think it is important to add these simulation exercises to the paper, we'll be happy to re-structure our simulation study accordingly.

Concerning the simulation exercises in (b), please also note the following: Our definitions of a correct and a spurious/false finding imply that the underlying target of interest is the unknown trend function m . This is different from the SiZer philosophy where the target is not the curve m itself but rather “the curve viewed at different resolution levels”. Formally, the target is the family of convolutions $m_h(u) = \int K(w)m(u + hw)dw$, that is, the family of smoothed versions m_h of the curve m . It would of course be possible to define correct and spurious/false findings in terms of the smoothed versions m_h . However, our main aim is to make rigorous confidence statements about the time intervals where the trend m (rather than the smoothed versions m_h) is increasing/decreasing. Hence, for us, it is more natural to think of the trend m itself as the underlying target.

Finally, many thanks for pointing us to Hannig et al. (2013). However, we have decided not to use the metrics of this paper in our simulation study for the following reason: The metrics defined in Hannig et al. (2013) can be used to measure how far the SiZer maps produced by different procedures (e.g. by dependent SiZer and our multiscale test) are from some oracle SiZer map. The oracle SiZer map is thus regarded as the target which should be approximated as well as possible by the procedures. This oracle SiZer map is based on a noiseless version of SiZer. (Specifically, it results from applying SiZer to the noiseless data $\{m(t/T) : t = 1, \dots, T\}$ rather than the noisy data $\{Y_{t,T} : t = 1, \dots, T\}$ and from plugging the true error

variance into the SiZer confidence intervals.) Hence, it is very natural to regard the oracle SiZer map as a benchmark for the map which SiZer produces from the actual (empirical or simulated) data. However, it is not clear to us why the oracle SiZer map should be a meaningful benchmark for the map produced by our multiscale test. Indeed, the oracle SiZer map cannot be regarded as a map which is generated by our multiscale test in an idealized noiseless situation (simply because our test approach is based on different test statistics and confidence intervals/critical values than SiZer). Hence, as far as we can see, it is not so meaningful to use the oracle SiZer map as a target when comparing our multiscale test with SiZer. Moreover, it would be a bit unfair to use it as a target since, by construction, it is closely related in structure to the SiZer test but not to our multiscale test.

(7) *Page 31, line 1-39: Can you plot the SiZer results on this data?*

We have applied both our multiscale test and row-wise dependent SiZer to the Central England temperature data. The SiZer maps produced by our test and dependent SiZer are plotted on p.33 of the revised paper.

References

- BERKES, I., LIU, W. and WU, W. B. (2014). Komlós-Major-Tusnády approximation under dependence. *Annals of Probability*, **42** 794–817.
- BROCKWELL, P. J. and DAVIS, R. A. (1991). *Time series: theory and methods*. New York, Springer.
- CHAUDHURI, P. and MARRON, J. S. (1999). SiZer for the exploration of structures in curves. *Journal of the American Statistical Association*, **94** 807–823.
- CHAUDHURI, P. and MARRON, J. S. (2000). Scale space view of curve estimation. *Annals of Statistics*, **28** 408–428.
- DONOHU, D. L. and JOHNSTONE, I. M. (1995). Adapting to unknown smoothness via Wavelet shrinkage. *Journal of the American Statistical Association*, **90** 1200–1224.
- DÜMBGEN, L. and SPOKOINY, V. G. (2001). Multiscale testing of qualitative hypotheses. *Annals of Statistics*, **29** 124–152.
- GRENNANDER, U. and SZEGÖ, G. (1958). *Toeplitz forms and their applications*. Cambridge University Press.
- HANNIG, J., LEE, T. C. and PARK, C. (2013). Metrics for SiZer map comparison. *Stat*, **2** 49–60.
- HANNIG, J. and MARRON, J. S. (2006). Advanced distribution theory for SiZer. *Journal of the American Statistical Association*, **101** 484–499.
- PARK, C., HANNIG, J. and KANG, K.-H. (2009). Improved SiZer for time series. *Statistica Sinica*, **19** 1511–1530.
- PARK, C., MARRON, J. S. and RONDONOTTI, V. (2004). Dependent SiZer: goodness-of-fit tests for time series models. *Journal of Applied Statistics*, **31** 999–1017.
- RONDONOTTI, V., MARRON, J. S. and PARK, C. (2007). SiZer for time series: a new approach to the analysis of trends. *Electronic Journal of Statistics*, **1** 268–289.
- WU, W. B. and POURAHMADI, M. (2009). Banding sample autocovariance matrices of stationary processes. *Statistica Sinica*, **19** 1755–1768.
- WU, W. B., WOODROOFE, M. and MENTZ, G. (2001). Isotonic regression: another look at the changepoint problem. *Biometrika*, **88** 793–804.
- XIAO, H. and WU, W. B. (2012). Covariance matrix estimation for stationary time series. *Annals of Statistics*, **40** 466–493.